

Roy S. Choi

323-430-1911 | roychoicareers@gmail.com | Open to relocate | [linkedin.com/in/roy-choi-9a839924a](https://www.linkedin.com/in/roy-choi-9a839924a) | github.com/IARPUS

EXPERIENCE

Software Engineer

Aug. 2025 – Present

Reevo — AI-native sales platform

\$80M raised — Kleiner Perkins & Khosla Ventures

- Designed and built an AI extraction pipeline (Temporal-orchestrated, Kafka event-driven) that converts meeting transcripts and emails into structured CRM updates end-to-end, with an idempotent, exactly-once write path and conflict resolution, batching multi-field updates into single bulk inserts — writing ~79–95K CRM field values/quarter at >99.6% reliability across ~1.2–2.1M reads/quarter.
- Built a post-meeting summarization feature with a claim-level citation system that grounds every generated statement to verbatim source spans, eliminating hallucinated attributions; adopted by ~290 reps with 15,210 views and 8,134 edits/quarter.
- Engineered the realtime collaborative-notes hot path on Cloudflare Workers (Y.js) to perform zero database operations per keystroke, with debounced, conditional history snapshots to control write amplification — sustaining sub-second multi-client sync across ~3,500 edits/quarter.
- Extended the “Ask Reevo” LLM assistant (~10K conversations/quarter across ~300 reps) onto Slack via a Slack integration that reuses the existing agent runtime across surfaces, letting reps query live CRM (customer relationship management) data without leaving Slack.
- Designed and shipped an automated pre-meeting prep feature that synthesizes CRM context (accounts, contacts, opportunities) into prep docs for sales calls, with a multi-source deep-research agent (concurrent source fan-out, partial-failure tolerance, synthesis under latency constraints) as one of its inputs.
- Diagnosed and eliminated recurring out-of-memory crashes in the prep workers by replacing a settings-polling loop with an event-driven model — removing the workflow state that Temporal’s stateful execution accrued from polling and cutting memory pressure ~90%.

Full-Stack Software Engineering Intern

May 2024 – Jul. 2024

Kintsugi Global

Los Angeles, CA

- Owned the manga catalog/homepage end-to-end — a browsable discovery surface with filtering, sorting, and pagination — in React and TypeScript over a Node.js/Express.js backend.
- Owned the in-page manga reader and integrated AWS (Amazon Web Services) S3 for scalable comic upload/retrieval; added Jest unit tests across components and routes and documented API endpoints in SwaggerHub.

PROJECTS

Bank Statement Analyzer | *React.js, FastAPI, PostgreSQL, OpenAI API*

May 2025

- Designed and built a solo full-stack app that parses US bank-statement PDFs (OpenAI API) and generates budgeting analysis and insights via LLMs.
- Built FastAPI microservices and an internal file system with PostgreSQL-backed metadata management for organized file retrieval.

Capstone — Siemens Digital | *React.js, Flask, Python, MongoDB*

Aug. 2024 – Dec. 2024

- On a team of 5, owned the endpoints and frontend for a Monte Carlo revenue-prediction simulation — Flask/Python REST routes with MongoDB persistence and interactive React visualizations; presented to Siemens VPs.

EDUCATION

University of Southern California

Aug. 2022 – May 2025

Bachelor of Science in Computer Science

Los Angeles, CA

TECHNICAL SKILLS

Languages: Python, C/C++, Java, JavaScript, TypeScript, SQL, HTML/CSS

Frameworks/Libraries: React.js, Next.js, Node.js, FastAPI, Flask, Django, Spring Boot, REST APIs, GraphQL, Pydantic, SQLAlchemy, Material UI, Tailwind, Jest, JUnit

Infrastructure: AWS (Amazon Web Services) — S3, Parameter Store; PostgreSQL, MongoDB, Temporal, Kafka, Cloudflare Workers, Docker

Developer Tools: Git, GitHub, Atlassian (Jira, Confluence), SwaggerHub